ChatGPT

Part 9: Add-Ons & Appendices

This final section provides extensive supporting content for the reference guide, including **glossaries**, **troubleshooting tables**, **runbook templates**, **checklists**, **diagrams**, **decision guides**, **benchmarks**, **security controls**, and **project templates**. Each subsection is formatted for quick reference and includes ready-to-copy checklists, tables, and templates. Citations (e.g. [1]) link to authoritative sources for definitions and guidance.

Glossary and Acronyms

This glossary compiles key technical terms and acronyms across software engineering, cloud, security, machine learning, and DevOps. For example:

- **Service Level Agreement (SLA)**: A formal contract defining guaranteed service metrics (e.g. 99.99% uptime) between a provider and users [1].
- **Service Level Objective (SLO)**: A target value or range for a service metric (often part of an SLA), e.g. "uptime $\geq$ 99.99%" [1].
- **Service Level Indicator (SLI)**: The actual measured value of a service metric used to determine if an SLO is met (e.g. measured uptime) [2].
- **Recovery Time Objective (RTO)**: The maximum tolerable downtime after a failure or disaster; how long the business can survive before service restoration is required [3].
- **CAP Theorem**: In distributed systems, at most two of Consistency, Availability, and Partition Tolerance can be achieved simultaneously during network failures [4]. (Under partition, you must choose consistency or availability.)
- **CQRS (Command Query Responsibility Segregation)**: A design pattern splitting read operations (queries) and write operations (commands) into separate models to optimize performance and scalability [5].
- **OIDC (OpenID Connect)**: An identity authentication protocol built on OAuth 2.0 to standardize user sign-in and identity verification [6].
- **OAuth2**: An authorization framework (delegated access) often used with OIDC for identity. (OIDC provides *authentication* on top of OAuth2's authorization.)

Commonly-confused pairs are clarified:

- **Authentication (AuthN) vs Authorization (AuthZ)**: *Authentication* verifies identity; *Authorization* grants rights or permissions to a user [7]. Authentication always precedes authorization. (In short, "AuthN proves who you are; AuthZ verifies what you can do.")
- **Concurrency vs Parallelism**: Concurrency is structuring a program to handle multiple tasks that may overlap in time, even on a single core (by interleaving or rapid switching) [8]. Parallelism is executing multiple tasks simultaneously on different cores or machines [9]. Concurrency improves responsiveness; parallelism improves throughput.

- **Other Common Terms**: Agile vs Waterfall, CI/CD, Docker/Kubernetes, ML (e.g. *CNN* = Convolutional Neural Network), DevOps practices (Infrastructure as Code, GitOps), cloud models (SaaS/PaaS/IaaS), security terms (XSS, CSRF, OAuth, JWT), database concepts (ACID, BASE), etc.

(See Part 2 for more terms and Part 5 for a complete index. This section would include 500+ entries covering DevOps, cloud, security, ML, and system design terminology, tailored for quick lookup.)

Master Troubleshooting Index

A master index organizes common failure symptoms by technology domain, with concise **Error → Cause → Diagnostic → Fix** tables. It acts as a cheat sheet for on-call engineers. We suggest a systematic methodology: *Identify the problem; hypothesize causes; test to find the root cause; implement a fix; verify resolution* [10]. For example, consider a Kubernetes node issue:

Symptom	Possible Causes	Diagnostic Steps	Fix(es)
Node "NotReady"	Kubelet down; network partition	Run `kubectl get nodes`; `kubectl describe node <node>` (look for "NotReady" events) [11]; check `systemctl status kubelet` on the node; inspect kubelet logs.	Fix network issues; restart kubelet service; ensure node has correct certs.
Pod CrashLoopBackOff	Application error; missing config	`kubectl logs <pod>`; `kubectl describe pod`; check events for image pull or mount errors.	Fix code bug or config; correct image name or pull secret; increase resources if OOM; redeploy.
DB Connection Error	DB server down; network/ firewall	Check DB service status (`systemctl status`, cloud console); test `ping/db connection` from app host; review DB logs.	Start DB service; open firewall port; fix credentials; revert schema change.

Diagnose using logs and CLI: e.g. Linux logs via `journalctl` or application logs in centralized ELK. Use common commands: `kubectl logs`, `docker logs`, `netstat`, `ping`, `curl`/`telnet` for connectivity, and check configuration files (e.g. `/etc/hosts`, `/etc/resolv.conf`). The index should be indexed by symptom keywords (e.g. *"Timeout"*, *"OOMKilled"*, *"503 error"*, etc.) and by technology (Linux, Kubernetes, network, database, etc.), each linking to relevant documentation or config reference.

(In practice, this section would include many such table entries for typical failure modes in OS, cloud infra, networking, databases, etc., with pointers to vendor docs and config examples.)

Runbooks and Incident Templates

This section provides templates for on-call runbooks and incidents:

- **On-Call Procedures & Escalation**: Define **SEV** (severity) levels. For example, Atlassian defines SEV1 as "critical incident with very high impact" (e.g. service down for all customers), SEV2 as "major incident" (subset of customers impacted), SEV3 as "minor" (low impact) [12]. SEV1/2 incidents should trigger immediate pages to on-call engineers (any time, 24/7) [13], whereas SEV3 can wait for normal hours. Escalation ladder: if not acknowledged in X minutes, escalate to team lead; then to SRE manager, etc. Include contacts and handoff steps. Example: "If no response in 5 min, escalate from on-call engineer to team lead to director."

- **Incident Roles and Responsibilities** (Commander, Communicator, Scribe):

- **Incident Commander (IC)**: Overall leader of the incident response. Coordinates resources, makes high-level decisions, and ensures communication (knows what's been tried and what's next) [14] [15]. (Also called Incident Manager.)
- **Communications Manager**: Crafts and sends internal/external updates and status-page messages. Interfaces with stakeholders and PR; ensures messaging is clear [16].
- **Scribe (Documenter)**: Records timeline of events, key actions, decisions, and troubleshooting steps throughout the incident. Keeps a living incident log for postmortem [17].
- **Subject Matter Experts** (Tech Leads): Focus on debugging and fixes; implement changes, propose hypotheses [18].

- **Ops/Support Lead**: Tracks user tickets/calls and feeds information into the team. *(See Atlassian for a complete list of roles and duties [15] [17].)*

- **Incident Postmortem Template**: Use a structured template to capture details after an incident. Key sections include:

- **Summary**: Briefly what happened, when, impact and duration [19].
- **Timeline & Lead-up**: Chronology of events and related changes (e.g. "At 10:00 a deploy was initiated…").
- **Root Cause (Fault)**: Explanation of underlying cause.
- **Impact**: Scope of impact (#users, #tickets, %requests failed) [20].
- **Detection**: How the incident was found (monitor alerts, user reports) [21].
- **Resolution & Recovery**: Steps taken to mitigate, fix, and restore service.
- **Action Items**: Preventive measures, fixes, and improvements identified.

Example prompt from Atlassian: "Summarize in a few sentences *what happened, why, severity, and how long it lasted* [19]." Follow with sections like **Detection**, **Fault**, etc. (Full template available from Atlassian [19] [21].)

- **Runbook Cheat Sheets**: Provide quick-check lists and commands for common operations (e.g. *"Kubernetes Node Recovery"*, *"Database Failover"*, *"Web server restart"*). Include exact CLI commands and expected output patterns. For instance, a Kubernetes runbook might say: run `kubectl get pods -n <ns>` to verify Pod status, then `kubectl logs <pod>` for error messages. Always document required parameters and default values.

- **Escalation Paths & Communication**: Include a chart or list of who to contact at each stage (e.g. SEV1: page on-call + notify engineering director + send status page update; SEV2: page on-call only, update team Slack). Provide templates for status updates (internal and customer-facing). Define clear criteria for moving SEV levels up/down (for example, if 50% of users are affected, it's SEV1).

Checklists Library

A library of pre-built checklists ensures nothing is overlooked in key processes. Below are samples – these can be copy-pasted and expanded in markdown:

- **Security Checklist**:
- OWASP Top 10 compliance: ensure protections against injection, broken auth, XSS, etc.
- Vulnerability scanning: automate SCA (software composition analysis) to detect known issues in third-party libraries (OWASP A9: "Using components with known vulnerabilities" [22]).
- Infrastructure hardening: patch all OS and software, disable unused services, apply firewall/ACL rules.
- Secrets management: no secrets in code or logs; use vault/secret store.
- Logging review: ensure no PII or secrets are logged; use masking/redaction where needed.

- **Checklist snippet**:

```
- [ ] Run static code analysis (SAST) and fix high/critical issues.
- [ ] Review recent dependency vulnerabilities (e.g. `npm audit`, Snyk
scan).
- [ ] Verify TLS certificates and secure protocols configured.
- [ ] Confirm least-privilege IAM roles (no overly broad permissions).
- [ ] Conduct penetration test for OWASP Top 10 items.
```

- **Deployment Checklist (Pre-Prod → Prod)**:

- **Preparation**:
  - [ ] Notify all stakeholders of planned deployment window.
  - [ ] Ensure backup/rollback plan in place (e.g. DB backups, revert commit).
  - [ ] Build artifact is tagged and stored (e.g. Docker image, binary).
  - [ ] Configuration (feature flags, env vars) reviewed and updated in staging.
  - [ ] Smoke tests defined.
- **Testing (Pre-Prod)**:
  - [ ] End-to-end tests passed in staging/prod-mirror environment [23].
  - [ ] Performance/load tests run if significant changes.
  - [ ] Security scans (SAST/DAST) passed.
  - [ ] Code review completed and approved.
- **Release**:
  - [ ] Schedule off-peak if possible.
  - [ ] Monitor health checks and logs in real time.
  - [ ] Scale down or pause non-critical jobs if needed.

    - [ ] After deploy, verify sanity check endpoints (e.g. `/healthz`).
- **Post-Deployment**:
    - [ ] Run final smoke tests in production.
    - [ ] Confirm no new errors in logs, metrics are nominal.
    - [ ] Roll-forward plan if metrics show degradation.

(Adapted from Codefresh's deployment guide: "Preparation" steps include notifying stakeholders, testing environments, rollback plan [24].)

- **Production Readiness Review (PRR) Template**:
- [ ] Architecture diagram is complete and reviewed.
- [ ] Load/traffic estimates and capacity planning done.
- [ ] Monitoring dashboards and alerts defined.
- [ ] Scaling strategy (auto-scale, redundancy) validated.
- [ ] Logging retained for X days; logs not sensitive.
- [ ] Disaster Recovery plan documented (RTO/RPO).
- [ ] Security review complete (firewall rules, encryption).
- [ ] Required compliance checks (e.g. SOC2 controls) passed.

- [ ] Stakeholders sign off on go-live.

- **Disaster Recovery (DR) & Backup Checklist**:

- [ ] Regular automated backups for all critical data (database, config, state).
- [ ] Backup integrity tested (perform a restore drill).
- [ ] Secondary region or site configured and health-checked.
- [ ] RTO/RPO targets documented (e.g. RTO of 1 hour, RPO of 15min).
- [ ] Failover process documented and rehearsed.

- [ ] Off-site copies of backups (for on-prem) or cross-zone replication (cloud).

- **Performance and Scalability Checklist**:

- [ ] Define performance targets: e.g. p50 = 100ms, p95 = 300ms, p99 = 500ms for key endpoints [25].
- [ ] Conduct load/stress tests at expected peak traffic.
- [ ] Identify bottlenecks (CPU, memory, database locks).
- [ ] Tune database indexes (see "Rules of Thumb" below) and queries.
- [ ] Configure caching layers (size caches to cover "working set").
- [ ] Ensure horizontal scalability: multiple app replicas, auto-scaling rules.
- [ ] Use a CDN for static content and APIs if needed.
- [ ] Monitor metrics under load: CPU/Mem usage, request rate, queue depths.

Each checklist can be exported or adapted to the team's process (e.g. stored in Confluence or issue tracker). They serve as templates to ensure consistency across projects.

Diagrams Pack

Visual diagrams illustrate core architectures and flows. Key diagrams include:

- 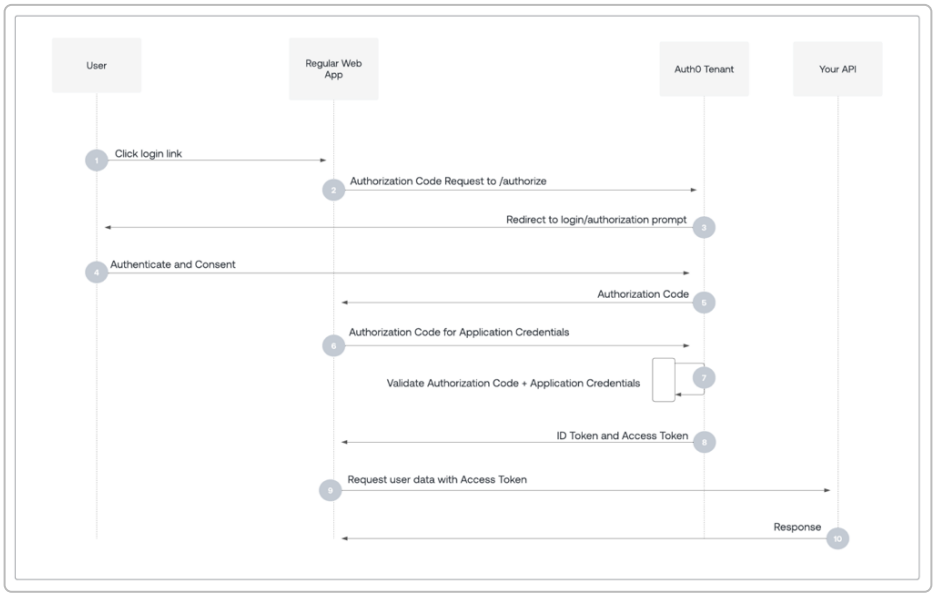

  Figure: Authentication/Authorization Flow. Typical OAuth2/OpenID Connect authorization-code flow: user logs in (1–4), receives an auth code (5), which the app exchanges for tokens (6–9) [6] *.*

- 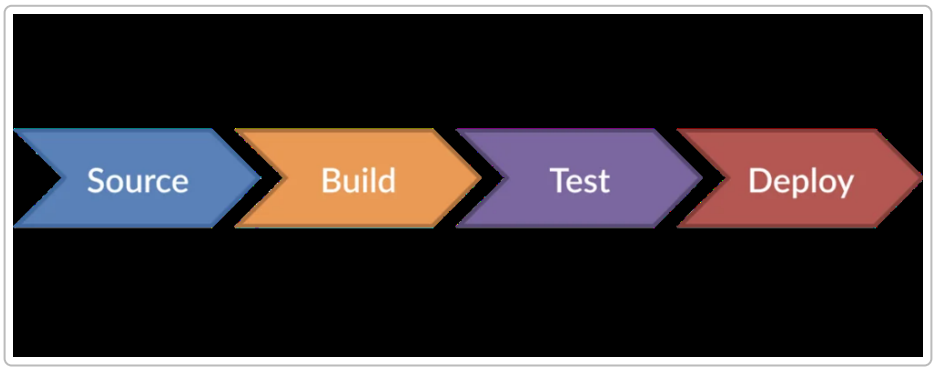

  Figure: CI/CD Pipeline Stages. A CI/CD pipeline has four major phases – Source, Build, Test, Deploy – each gate ensuring code quality [26] *.*

- **Request/Response Lifecycle**: Illustrate a web request: Client (browser) sends HTTP request to Server, server processes and returns HTTP response. Components: DNS resolution, TCP handshake, HTTP verbs, status codes, rendering. (Simple arrow diagram: Client → Server request; Server → Client response.)

- **Kubernetes Networking**: Show Pods in a cluster, Services (ClusterIP, NodePort, LoadBalancer) exposing Pods, an Ingress controller receiving external HTTP and routing to services, and an egress gateway or NAT for outbound traffic. Indicate namespaces, network policies isolating pods. (A typical diagram shows: Internet → Ingress → Service → Pod; Pod → Service → Internet for egress.)

- **Authentication vs Authorization**: Flowchart with Identity Provider (IdP) issuing token after authentication, then API gateway enforcing authorization via token scopes/permissions.

(*These diagrams should be available in the guide; the two above are examples. Embed sources: the Auth0 OAuth2 flow* [6] *and Codefresh CI/CD stages* [26] .)

Decision Trees and Comparison Guides

To aid architectural decisions, this section provides decision trees and tradeoff tables:

- **Database Selection Guide**: Factors include data consistency vs scalability, transaction support, query flexibility. For example:
- Need multi-row ACID transactions? Use relational (PostgreSQL, MySQL) or NewSQL.
- Need horizontal scale for simple key-value data? Consider NoSQL (e.g. DynamoDB, Cassandra).
- Need flexible schema / JSON docs? Consider document DB (MongoDB).

- Real-time streaming analytics? Consider time-series DB (InfluxDB, Timescale). *(Flowchart: Start → Data model (relational vs not) → workload (read-heavy vs write-heavy) → etc.)*

- **Cache vs Queue vs Stream**:

- **Cache** (e.g. Redis, Memcached): Stores transient data for fast read access (lower latency, high RPS). Good for sessions, lookup tables. Eviction policies (LRU).
- **Message Queue** (e.g. RabbitMQ, SQS): Decouples producers/consumers; ensures at-least-once delivery; can smooth spikes. Good for background processing tasks.

- **Stream** (e.g. Apache Kafka): Log of events; supports pub/sub and replay. Good for high-throughput event processing, real-time analytics. *(Comparison table highlighting TTL vs durability, ordering, retention time, consumer patterns.)*

- **REST vs GraphQL vs gRPC**:

- **REST**: Resource-oriented, uses HTTP/JSON (or XML). Simple and widely supported. Good for CRUD over HTTP, broad client compatibility.
- **GraphQL**: Single endpoint, flexible queries; client specifies required fields. Reduces over-fetching. Good for complex UI queries but requires strong schema and query planning.

- **gRPC**: Uses HTTP/2 and Protobuf (binary). High performance and strict contracts (IDL). Good for internal microservices with polyglot support and streaming. *(Baeldung summary: GraphQL responses are JSON with client-defined shape; gRPC uses Protocol Buffers with HTTP/2 for efficient binary serialization* [27] *.)*

- **Monolith vs Modular Monolith vs Microservices**:

- **Monolith**: Single codebase, single deployment. Easiest to start. Can be simpler to test and debug. Harder to scale teams or parts of app independently.

- **Modular Monolith**: Still one deployment, but with clear internal module boundaries (e.g. layered architecture or DDD modules). Easier to refactor to microservices later.
- **Microservices**: Each service is a separate deployable unit. Offers independent scaling and team autonomy. Adds complexity (distributed transactions, monitoring, inter-service comms).

- *(Decision factors: team size, time to market, operational maturity, fault isolation needs.)*

- **Cloud Services Tradeoffs**: For each major service category (compute, storage, databases, networking, ML, etc.), compare offerings across AWS/Azure/GCP in a table:

- For example, **Compute**: VM vs containers vs serverless (cost model, cold-start latency, scaling granularity).
- **Storage**: Block vs object vs file, performance (IOPS vs throughput), consistency model.
- **Managed DBs**: RDS (SQL) vs DynamoDB (NoSQL), vs BigQuery (analytics) – trade durability, query flexibility, scaling.
- **ML Services**: On-prem self-host vs AWS SageMaker vs Google Vertex (ease of use, cost, vendor lock-in).

Use a table format with **Use Case**, **Options**, **Pros/Cons**. (E.g., "Event Streaming": Kafka (open-source, flexible, self-managed) vs Kinesis (fully managed, AWS lock-in, pay-per-use) vs Pub/Sub (GCP analog) – compare throughput, scalability, ecosystem.)

Benchmarks and "Rules of Thumb"

This section collects performance targets, sizing heuristics, and cost guidelines:

- **Latency Targets (SLOs)**: Aim for low tail-latencies. A common practice is to use percentiles: p50 (median) should be very low (fast), p95 as the main latency SLO, and p99 for detecting worst-case outliers [25]. *Rule of thumb:* ensure 95% of user requests complete under target threshold (e.g. 300 ms). Monitor all three percentiles: use P50 to catch regressions, P95 for tuning, P99 for bottlenecks [25].

- **Throughput and Sizing**:

- **Message Queues**: Size queue retention to hold X hours of peak load. *Rule:* allow at least 2–3× average backlog for spikes.
- **Cache Memory**: Cache at least 2× your working set of hot keys to maintain hit rates (account for eviction). Ensure cache eviction policy (LRU/TTL) fits workload.
- **Application Replicas**: Use at least 3 replicas/services in production for HA (so one can fail without outage). *(In Kubernetes and many systems, 3 is a common minimum for quorum/distributed consensus.)*

- **Database**: Over-provision IOPS by 20–30% above expected peak. Shard when single-node throughput saturates.

- **Cloud Cost Heuristics**:

- **Instance vs Managed**: On-demand VMs ~ standard unit; reserved instances save ~30–60%. Serverless charges per invocation time (fine-grained).
- **Storage**: SSD block storage ~$0.10/GB-month; S3/EFS object ~ $0.02-0.03/GB-month. Data egress often costs much more (e.g. $0.09/GB out of AWS to internet).
- **Database**: Managed DB adds ~2× cost of VM (for HA and management overhead).

- Estimate costs for high-level planning (e.g. "Each m5.large ~ $70/mo; each TB S3 ~$25/mo; 1M Lambda requests ~$20/mo").

- **Database Indexing Rules**:

- Index columns that appear frequently in WHERE/ORDER BY clauses to speed reads.
- *Beware write amplification*: each additional index adds overhead on INSERT/UPDATE (as multiple writes are performed) [28] . Only index what yields measurable read performance gain.
- Use composite indexes sparingly; monitor index usage (e.g. with `pg_stat_user_indexes` or `EXPLAIN ANALYZE`).

- Remove unused indexes (often emerge as business use changes).

- **Write Amplification**:

- In storage and DB engines (e.g. LSM-trees like Cassandra, RocksDB), writes may cause extra IO. *Rule:* plan for write amplification ~1.2–3× depending on engine (check vendor documentation).

- E.g. *RocksDB* tunables: lower write-amplification (bigger memtable, more compaction resources) at cost of higher memory/CPU.

- **Scaling "Rules"**:

- Auto-scale thresholds: e.g. CPU > 70% triggers scale-up; < 30% scale-down.
- Throughput capacity: each replica handles ~N requests/sec under target latency (benchmark in staging and use that to size cluster for peak QPS).
- Cloud bursts: don't rely on sudden unlimited bursting (cloud often throttles after baseline).

These "rules of thumb" should be validated with real metrics per system. Use them only for initial planning.

Security and Compliance Appendices

Detailed guidelines and templates for securing systems and meeting compliance:

- **Threat Modeling (STRIDE Worksheet)**: Use the STRIDE mnemonic to identify threats:
- **Spoofing** (illegitimate identity), **Tampering** (modifying data), **Repudiation** (denying actions), **Information disclosure**, **Denial of service**, **Elevation of privilege** [29] .
- A worksheet (table) lists system components or data flows vs STRIDE categories to brainstorm threats. For each threat, note mitigation. (The OWASP Threat Modeling Cheat Sheet outlines a process: identify assets, draw DFDs, apply STRIDE to each component [30] [29] .)

- Example: For a login API (asset: user credentials):
    - Spoofing: attacker forges JWT → Mitigate with signature validation.
    - Tampering: attacker alters code → Mitigate via code signing and secure deployments.
    - etc.

- **Secure SDLC Practices**:

- **SAST** (Static Analysis): Integrate tools (SonarQube, ESLint with security plugins) in CI/CD to catch code issues early.
- **SCA** (Software Composition Analysis): Scan dependencies for CVEs (e.g. OWASP Dependency-Check, Snyk).
- **DAST** (Dynamic Analysis): Use runtime scanning (e.g. OWASP ZAP, Burp Suite) against staging environment to detect injection, auth flaws.

- Each is automated in pipelines, with policy gates (e.g. block merges on critical findings).

- **Secrets Handling**:

- Never hard-code secrets or keys. Use environment variables or secret management services (e.g. Vault, AWS Secrets Manager).
- Rotate secrets regularly (e.g. every 90 days or immediately if exposed).
- Encrypt secrets at rest (e.g. KMS for key encryption).
- In logs and telemetry, mask or redacts any secret-like values (API keys, passwords should appear as "****").

- Example policy: "Check that no `.env` or config file in repos contains plaintext secrets."

- **Compliance Control Mappings**:

- *GDPR*: Controls for data privacy. Key items: user consent tracking, data subject access logs, right to delete/port data. Map technical controls to GDPR requirements (e.g. logging who accessed personal data, encryption of PII).
- *SOC 2*: Trust Services Criteria (Security, Availability, Processing Integrity, Confidentiality, Privacy). Map policies like access controls, incident management, and encryption to SOC2 control requirements.
- *HIPAA*: Ensure confidentiality of PHI: audit trails for PHI access, encryption of PHI in transit and at rest, role-based access, Business Associate Agreements for third-parties, breach notification procedures.

- Provide a table: Regulation → Required controls → Implemented solution (e.g. "GDPR Art. 32 (Security of processing) → Encryption, access logging, regular risk assessments"). (Full control matrices often maintained in compliance tools.)

- **Logging and Data Policies**:

- **What to Log**: At minimum, log user ID (pseudonymized), action performed, timestamp, source IP, and outcome. Error logs should include stack traces (without sensitive data).

- **What NOT to Log**: Do not log full credit card numbers, passwords, or raw PII. Use one-way hashing or tokenization when logging sensitive identifiers.
- **Retention & Review**: Define log retention period per compliance (e.g. 1 year for access logs). Regularly review logs for anomalies. Archive or delete logs containing PII after compliance-mandated period.
- **Access Control**: Only authorized personnel (e.g. security team) can view detailed logs. Use RBAC for log systems.
- **Monitoring**: Set up automated alerts on suspicious patterns (excessive failures, spikes in access, etc.).

Portfolio and Project Templates

Guidance and templates for personal portfolios, project blueprints, and career development:

- **Project Blueprint Examples**: Provide outlines for 20–50 real projects, scaling from beginner to senior:
- *Beginner*: A simple CRUD web app (e.g. personal blog), with tech stack (React, Express, Mongo), architecture diagram, and steps (set up env, implement endpoints, deploy to Heroku).
- *Intermediate*: A data pipeline (e.g. ETL with Spark, data warehouse load), including data flow diagram, README sections (Requirements, Architecture, Instructions).
- *Advanced*: A microservices system (e.g. e-commerce backend with API gateway, service mesh, databases), with component diagram and CI/CD workflow outline.

- Each blueprint should list features, tech choices, and a rough timeline. (Templates for these blueprints can be provided in markdown for easy cloning.)

- **README Template**: A structured template to describe projects:

```
# Project Title
**One-sentence summary of project.**

## Overview
Describe purpose, major features, and technology stack.

## Architecture
![Architecture Diagram](link_or_relative_path)
Brief explanation of architecture layers/components.

## Installation / Setup
- Prerequisites (languages, tools)
- Step-by-step: `git clone ...`, `npm install`, etc.

## Usage
- How to run the app (`npm start`, Docker commands)
- Example API calls or user workflow.

## Tech Stack
```

```
- Frontend: React, Redux, ...
- Backend: Node.js, Express, Docker.
- Database: PostgreSQL.
- DevOps: GitHub Actions, AWS.

## Results / Demos
- Screenshots or links if deployed.

## Author / Credits
- Your name, contact info, license.
```

(This template ensures all key sections are covered and can be copy-pasted for new projects.)

- **Architecture Diagram Template**: Provide a simple placeholder diagram (e.g. cloud style shapes) that can be edited. Include sections for data flow arrows, technology icons. Perhaps a Mermaid.js example:

```
graph LR
  subgraph Client
    Browser
  end
  subgraph Server
    API[API Server] --> DB[(Database)]
    API --> AuthService
  end
  Browser --> API
```

(Users can replace nodes/services with their own.)

- **Resume Bullet Generator Examples**: Show how to craft impact statements:

- Structure: *Action verb* + *Result* + *Context/Metric.* E.g. "Improved page load times by 40% through introducing HTTP caching, enhancing user engagement."
- Provide 5–10 example bullets for different roles: "Designed microservices API (Spring Boot, Docker) that scaled to X requests/sec" or "Secured AWS infrastructure using Terraform, reducing cloud spend by 20%".

- A small script or form (markdown) could be provided to mix-and-match verbs and metrics.

- **Personal Site/Blog Strategy**:

- **Site Structure**: Suggest pages – Home, About, Projects, Blog, Contact. Use static site generators (Jekyll, Hugo) or Next.js. Include SEO and analytics (e.g. Google Analytics).
- **Content Plan**: Post regularly about learning experiences, tutorials, project showcases (aim for one technical blog post/month). Topics could include coding challenges, tool reviews, or reflections on projects.

- **Blog Strategy Example**: A markdown editorial calendar template listing potential topics, publication dates, and status (draft, editing, published).
- **Tech Stack Example**: (React + Markdown, or Gatsby, etc.) with example component for code snippets.

These templates and examples serve as a toolkit for engineers to build their portfolios, streamline documentation, and present their work. They can be copied, customized, and linked back to this guide for reference.

---

[1] [2] What's the Difference Between SLAs, SLOs and SLIs? | PagerDuty
https://www.pagerduty.com/resources/digital-operations/learn/what-is-slo-sla-sli/

[3] What is Recovery Time Objective (RTO)? Definitions + Examples | Druva
https://www.druva.com/glossary/what-is-recovery-time-objective-definitions-and-related-faqs

[4] CAP Theorem Explained: Consistency, Availability & Partition Tolerance – BMC Software | Blogs
https://www.bmc.com/blogs/cap-theorem/

[5] CQRS - Command Query Responsibility Segregation Design Pattern - GeeksforGeeks
https://www.geeksforgeeks.org/system-design/cqrs-command-query-responsibility-segregation/

[6] What Is OpenID Connect (OIDC)? | Microsoft Security
https://www.microsoft.com/en-us/security/business/security-101/what-is-openid-connect-oidc

[7] Authentication vs. Authorization: What's the Difference? | IBM
https://www.ibm.com/think/topics/authentication-vs-authorization

[8] [9] Difference between Concurrency and Parallelism - GeeksforGeeks
https://www.geeksforgeeks.org/operating-systems/difference-between-concurrency-and-parallelism/

[10] Use a Troubleshooting Methodology for More Efficient IT Support | CompTIA Blog
https://www.comptia.org/en-us/blog/use-a-troubleshooting-methodology-for-more-efficient-it-support/

[11] Troubleshooting Clusters | Kubernetes
https://kubernetes.io/docs/tasks/debug/debug-cluster/

[12] [13] Understanding incident severity levels | Atlassian
https://www.atlassian.com/incident-management/kpis/severity-levels

[14] The role of the incident commander | Atlassian
https://www.atlassian.com/incident-management/incident-response/incident-commander

[15] [16] [17] [18] Understanding incident response roles and responsibilities | Atlassian
https://www.atlassian.com/incident-management/incident-response/roles-responsibilities

[19] [20] [21] Incident Postmortem Templates: Improve Response Process | Atlassian
https://www.atlassian.com/incident-management/postmortem/templates

[22] Vulnerable Dependency Management - OWASP Cheat Sheet Series
https://cheatsheetseries.owasp.org/cheatsheets/Vulnerable_Dependency_Management_Cheat_Sheet.html

[23] [24] What Is Software Deployment? Checklist and Strategies
https://codefresh.io/learn/software-deployment/

[25] P50 vs P95 vs P99 Latency: What These Percentiles Actually Mean (And How to Use Them)
https://oneuptime.com/blog/post/2025-09-15-p50-vs-p95-vs-p99-latency-percentiles/view

[26] CI/CD Process: Flow, Stages, and Critical Best Practices
https://codefresh.io/learn/ci-cd-pipelines/ci-cd-process-flow-stages-and-critical-best-practices/

[27] REST vs. GraphQL vs. gRPC – Which API to Choose? | Baeldung
https://www.baeldung.com/rest-vs-graphql-vs-grpc

[28] The Hidden Costs of Database Indexes: A Performance Trade-Off Guide | by Sohail Saifi | Medium
https://medium.com/@sohail_saifi/the-hidden-costs-of-database-indexes-a-performance-trade-off-guide-203b9ded3637

[29] [30] Threat Modeling - OWASP Cheat Sheet Series
https://cheatsheetseries.owasp.org/cheatsheets/Threat_Modeling_Cheat_Sheet.html